

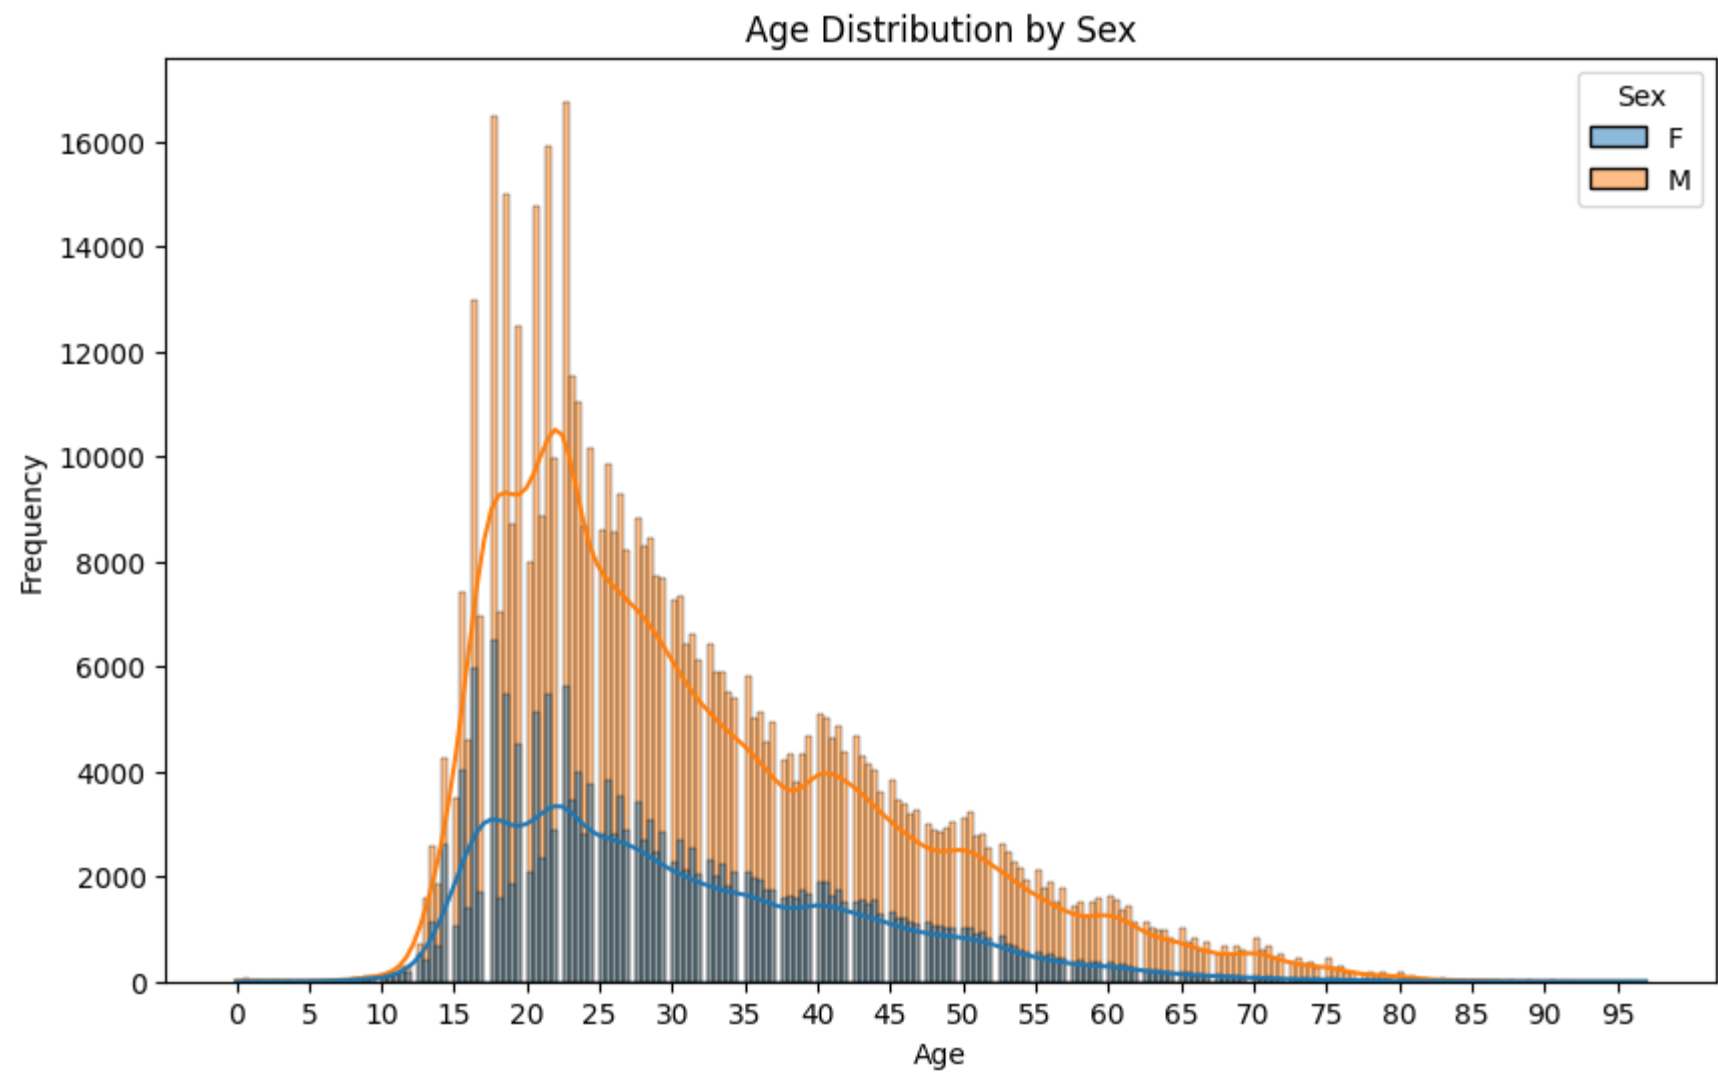
```
In [ ]: import pandas as pd
import gdown
import gzip
from io import BytesIO
import copy
# Here we load the dataset into a Pandas DataFrame using Pandas
# Note for some reason I could no longer upload the file containing the dataset so I decided to host the file on google drive.
file_id = '1g_B-ms_dmZQWzuRLOBJveFN6b7sg9AJ7'
url = f'https://drive.google.com/uc?id={file_id}'
output = 'data.csv.gz'
# The reason to us gzip is that the dataset is compressed because it is on the larger size at about 57MB
gdown.download(url, output, quiet=True)
with gzip.open(output, 'rt') as file:
# the file is a csv file
    df = pd.read_csv(file)
# We need to copy the df for our secondary cleaning and data transformation.
df_copy = copy.deepcopy(df)
```

```
<ipython-input-1-183a18b708db>:15: DtypeWarning: Columns (35) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv(file)
```

## Section 1

In this section we perform a preliminary data exploration of the dataset which can be downloaded from Kaggle → [Dataset](#)

```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt
# we create a temporary copy of the df to display some charts
temp = copy.deepcopy(df)
temp['Sex'] = temp['Sex'].replace({0: 'Female', 1: 'Male'})
# The folloing graph prints the distrubution of age by the sex of the lifters, this helps visualize both sexes at the same time
plt.figure(figsize=(10, 6))
sns.histplot(data=temp, x='Age', hue='Sex', kde=True)
plt.title('Age Distribution by Sex')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.xticks(range(0, 100, 5))
plt.show()
# df.head gives a an overview of what the raw data looks like we have 37 distinct columns, most of which we will find unsuable
print(df.head())
```




	Name	Sex	Event	Equipment	Age	AgeClass	Division	BodyweightKg	\
0	Abbie Murphy	F	SBD	Wraps	29.0	24-34	F-OR	59.8	
1	Abbie Tuong	F	SBD	Wraps	29.0	24-34	F-OR	58.5	
2	Ainslee Hooper	F	B	Raw	40.0	40-44	F-OR	55.4	
3	Amy Moldenhauer	F	SBD	Wraps	23.0	20-23	F-OR	60.0	
4	Andrea Rowan	F	SBD	Wraps	45.0	45-49	F-OR	104.0	

	WeightClassKg	Squat1Kg	...	McCulloch	Glossbrenner	IPFPoints	Tested	\
0	60	80.0	...	324.16	286.42	511.15	NaN	
1	60	100.0	...	378.07	334.16	595.65	NaN	
2	56	NaN	...	38.56	34.12	313.97	NaN	
3	60	-105.0	...	345.61	305.37	547.04	NaN	
4	110	120.0	...	338.91	274.56	550.08	NaN	

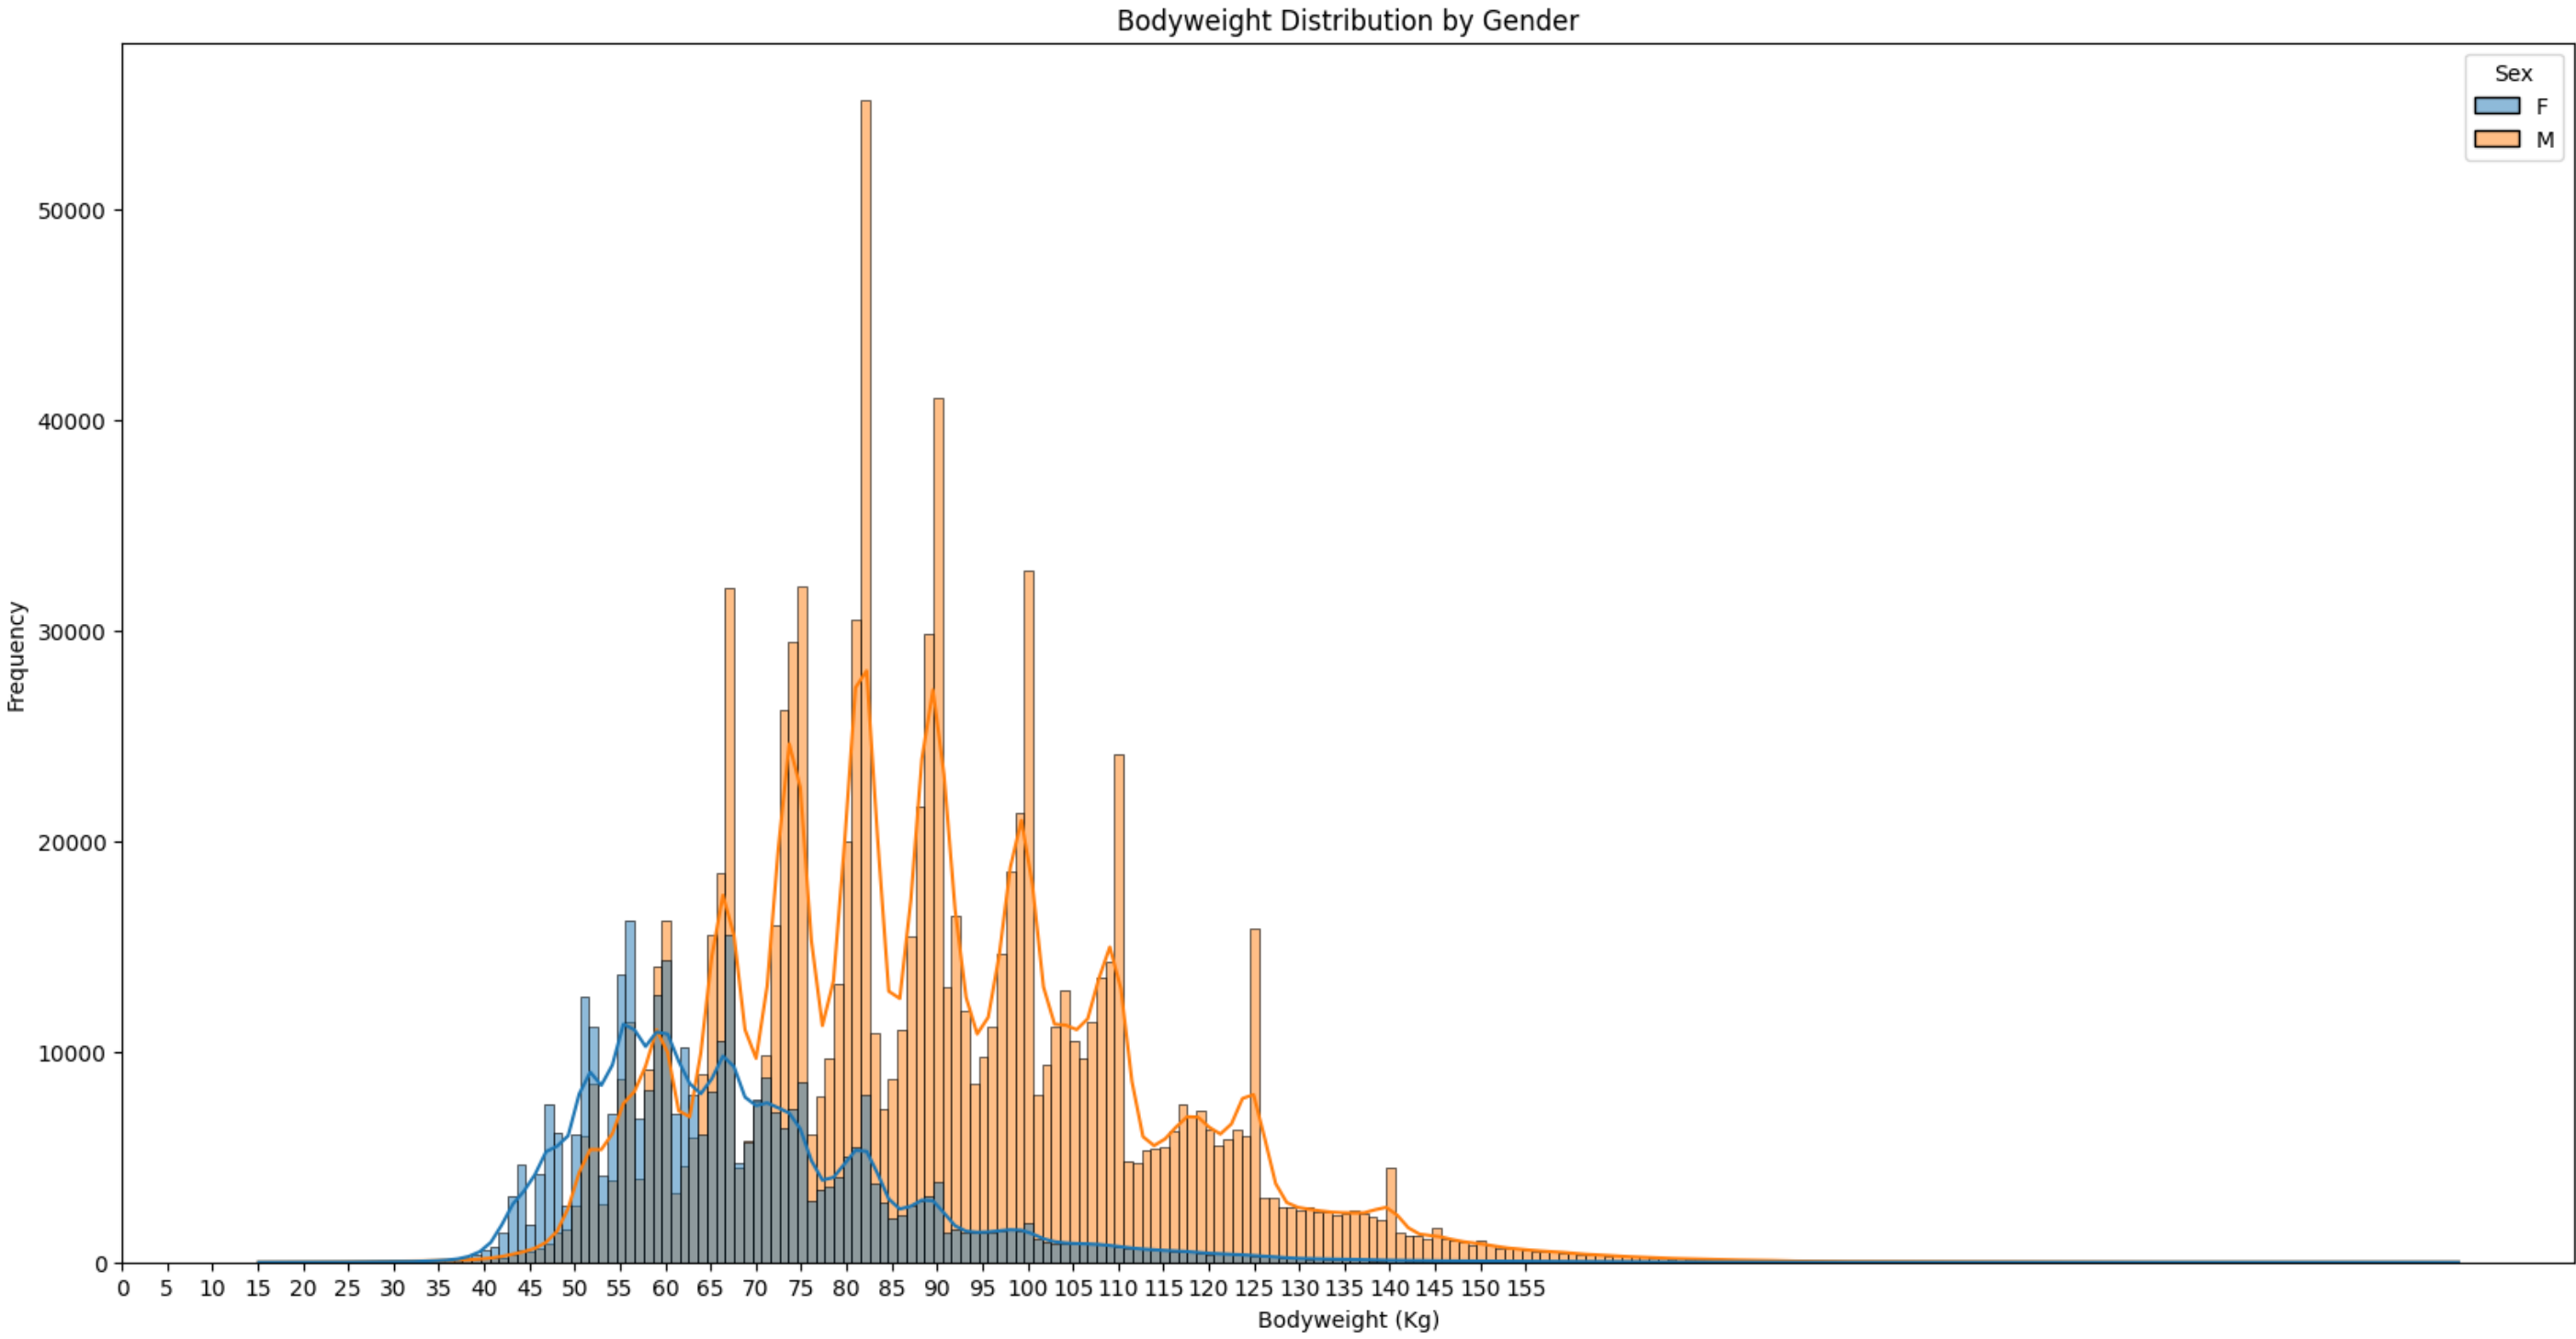
	Country	Federation	Date	MeetCountry	MeetState	MeetName
0	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
1	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
2	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
3	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
4	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup

[5 rows x 37 columns]

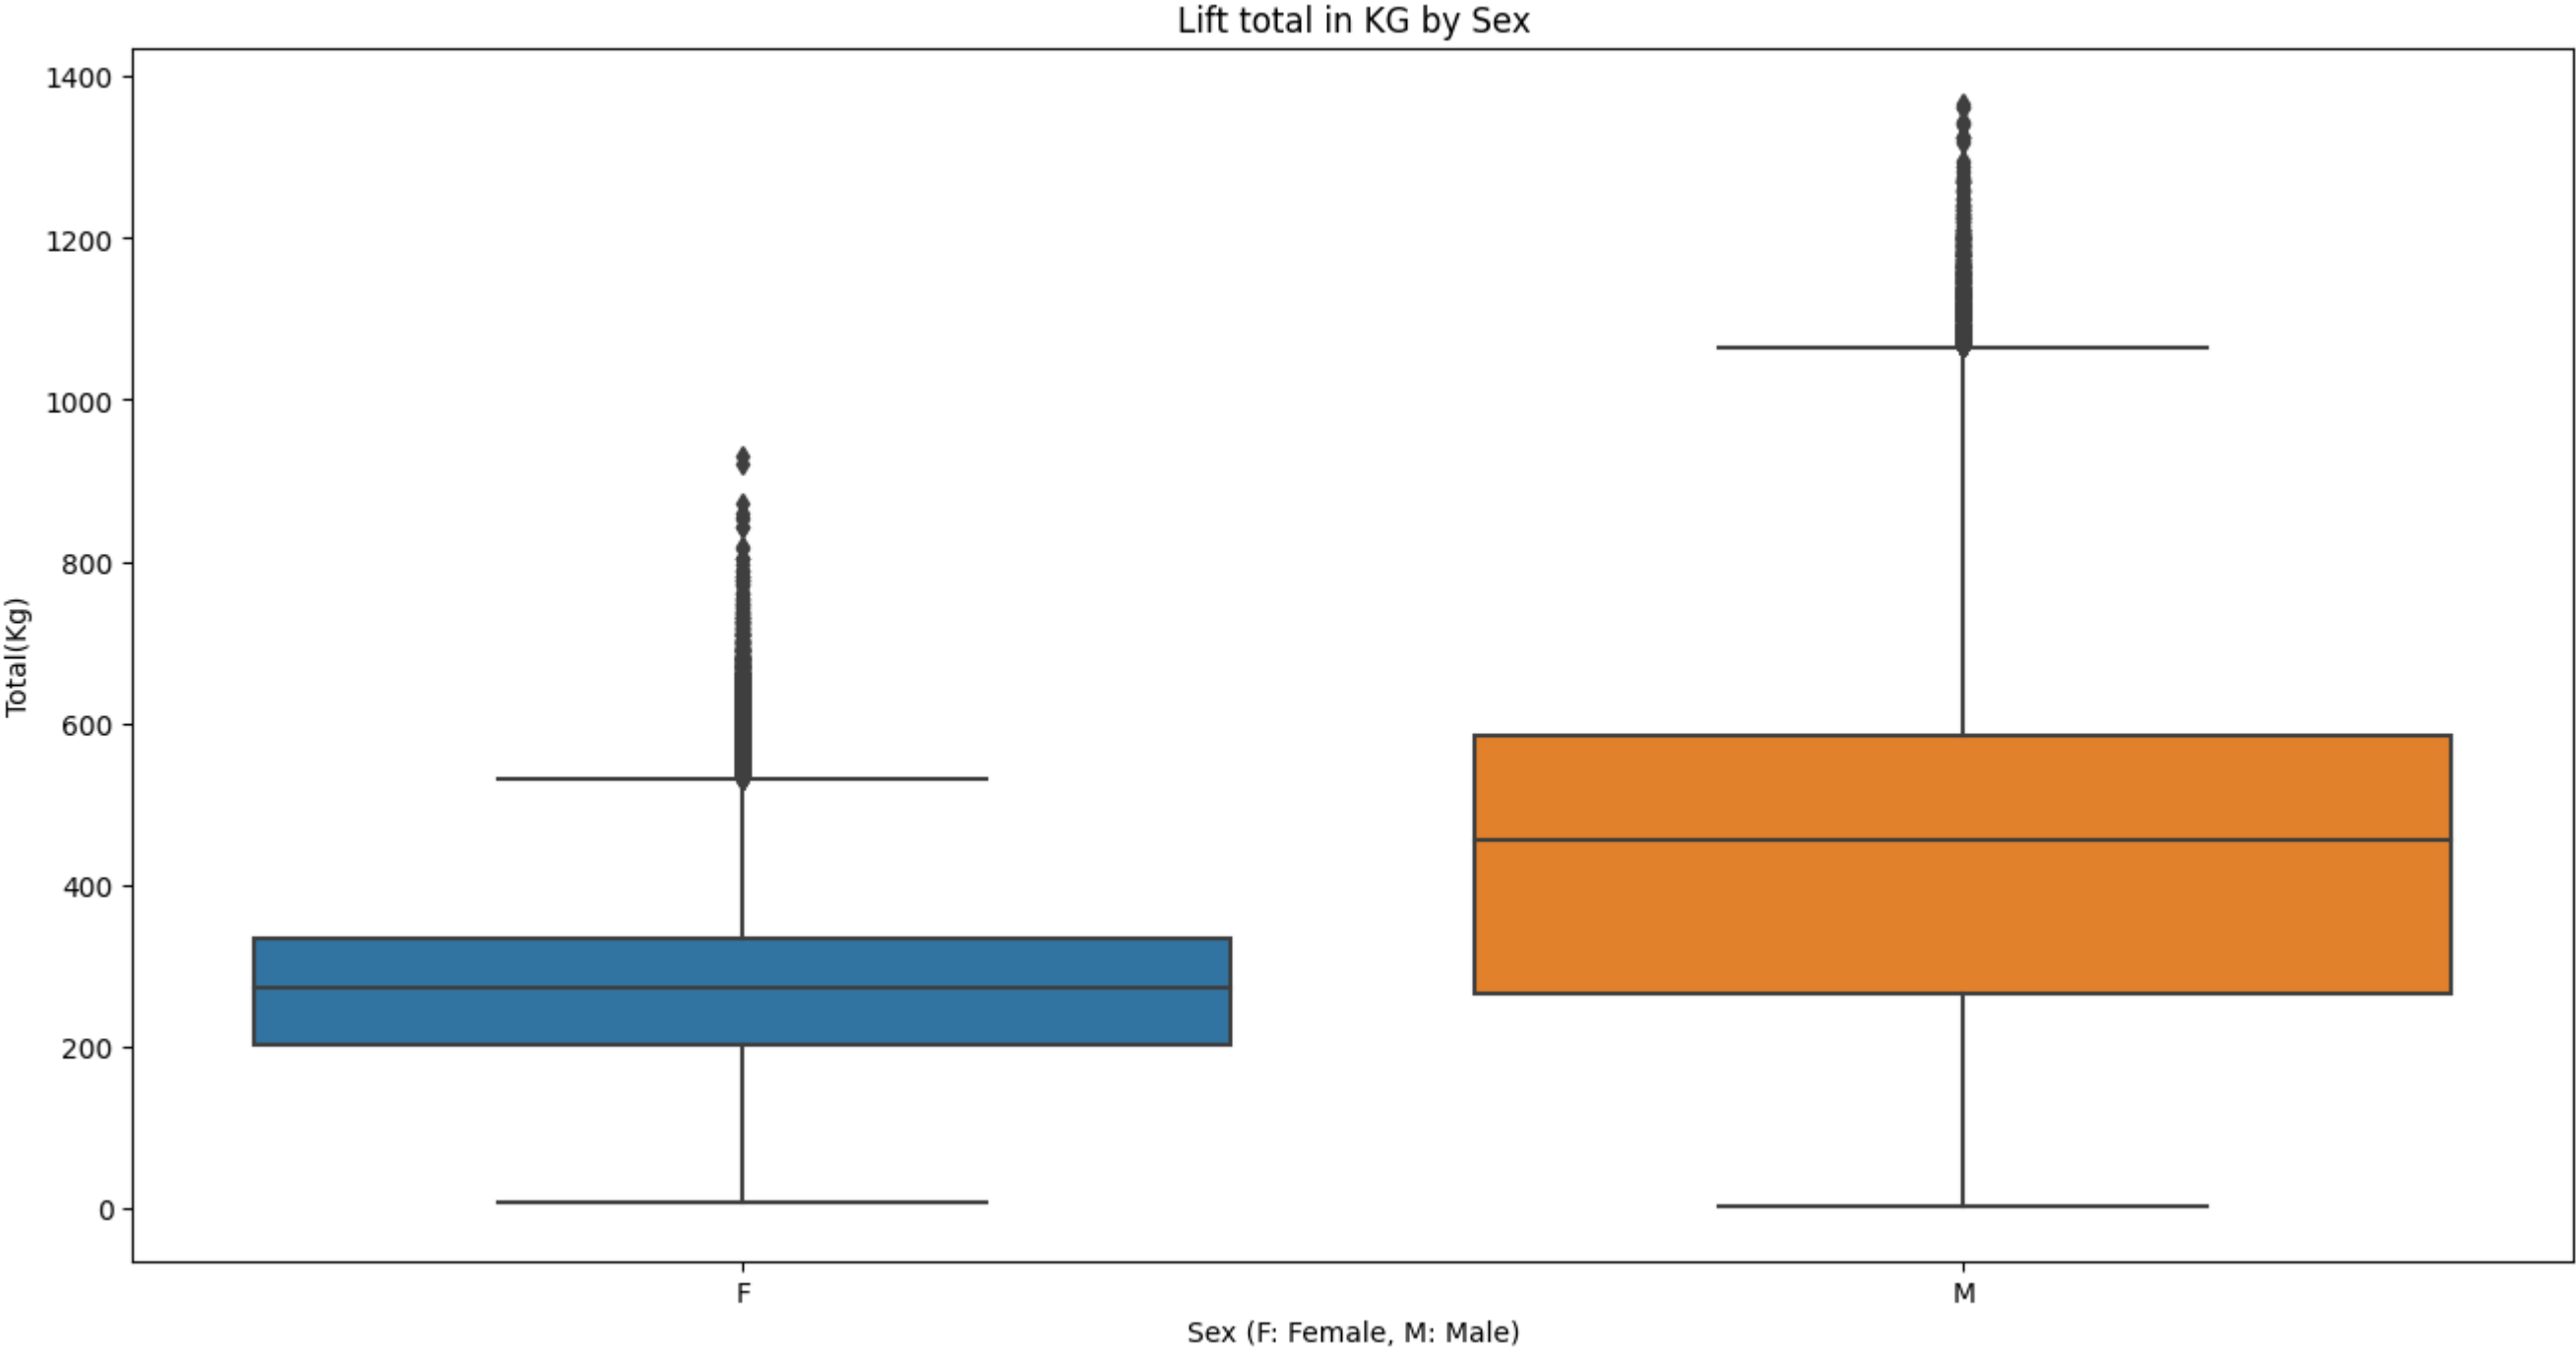
```
In [ ]: # This graph allows us to see how many male and female competitors are in the raw data
plt.figure(figsize=(15, 7))
axis = sns.countplot(data=df, x='Sex')
plt.title('Sex Distribution')
plt.xlabel('Sex (F: Female, M: Male)')
plt.ylabel('Count')
axis.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, loc: "{:,}".format(int(x))))
plt.show()
```



```
In [ ]: # This graph visulises the body weight distribution of the lifters by gender
plt.figure(figsize=(20, 10))
sns.histplot(data=df, x='BodyweightKg', hue='Sex', kde=True, discrete=True)
plt.title('Bodyweight Distribution by Gender')
plt.xlabel('Bodyweight (Kg)')
plt.ylabel('Frequency')
plt.xticks(range(0, 160, 5))
plt.show()
```



```
In [ ]: # The graph shows a the distribution of total weight lifted by competitors, most competitors are in the range 200 to 600 kgs, the outliers are due to equipped lifters which we be
# Since we care about raw powerlifting, most of the
plt.figure(figsize=(16, 8))
sns.boxplot(data=df, x='Sex', y='TotalKg')
plt.title('Lift total in KG by Sex')
plt.xlabel('Sex (F: Female, M: Male)')
plt.ylabel('Total(Kg)')
plt.show()
```



```
In [ ]: # This section provides some statistics about the raw data, as we can see we have many duplciated data and missing values
data_description = df.describe()
print(data_description)
correlation_matrix = df.corr()
print(correlation_matrix)
missing_values = df.isnull().sum()
print("Missing values count: "+str(missing_values))
total_missing_values = df.isna().sum().sum()
print("Total NaN values: "+ str(total_missing_values))
duplicated_rows_count = df.duplicated().sum()
print("Total duplicated rows count: "+str(duplicated_rows_count))
print(df.head())
```

	Age	BodyweightKg	Squat1Kg	Squat2Kg	\
count	757527.000000	1.406622e+06	337580.000000	333349.000000	
mean	31.501570	8.422503e+01	114.102442	92.155846	
std	13.371707	2.322011e+01	147.143021	173.701524	
min	0.000000	1.510000e+01	-555.000000	-580.000000	
25%	21.000000	6.670000e+01	90.000000	68.000000	
50%	28.000000	8.180000e+01	147.500000	145.000000	
75%	40.000000	9.915000e+01	200.000000	205.000000	
max	97.000000	2.580000e+02	555.000000	566.990000	
	Squat3Kg	Squat4Kg	Best3SquatKg	Bench1Kg	Bench2Kg \
count	323842.000000	3696.000000	1.031450e+06	499779.000000	493486.000000
mean	30.056842	71.356870	1.740049e+02	83.892373	55.065745
std	200.413385	194.522045	6.923931e+01	105.196350	130.302229
min	-600.500000	-550.000000	-4.775000e+02	-480.000000	-507.500000
25%	-167.500000	-107.840000	1.224700e+02	57.500000	-52.500000
50%	110.000000	135.000000	1.678300e+02	105.000000	95.000000
75%	192.500000	205.000000	2.175000e+02	145.000000	145.000000
max	560.000000	505.500000	5.750000e+02	467.500000	487.500000
	Bench3Kg	...	Deadlift1Kg	Deadlift2Kg	Deadlift3Kg \
count	478485.000000	...	363544.000000	356023.000000	339947.000000
mean	-18.520481	...	162.700840	130.228378	12.995484
std	144.225726	...	108.681438	162.680134	215.052488
min	-575.000000	...	-461.000000	-470.000000	-587.500000
25%	-140.000000	...	125.000000	115.000000	-210.000000
50%	-60.000000	...	180.000000	177.500000	117.500000
75%	117.500000	...	226.800000	230.000000	205.000000
max	478.540000	...	450.000000	460.400000	457.500000
	Deadlift4Kg	Best3DeadliftKg	TotalKg	Wilks	McCulloch \
count	9246.000000	1.081808e+06	1.313184e+06	1.304407e+06	1.304254e+06
mean	78.914945	1.872585e+02	3.956148e+02	2.882247e+02	2.960682e+02
std	192.605159	6.232821e+01	2.011420e+02	1.231805e+02	1.249700e+02
min	-461.000000	-4.100000e+02	2.500000e+00	1.470000e+00	1.470000e+00
25%	-110.000000	1.383500e+02	2.325000e+02	1.979000e+02	2.048200e+02
50%	145.150000	1.850000e+02	3.787500e+02	3.052000e+02	3.120300e+02
75%	210.000000	2.300000e+02	5.400000e+02	3.745600e+02	3.837600e+02
max	418.000000	5.850000e+02	1.367500e+03	7.793800e+02	8.044000e+02
	Glossbrenner	IPFPoints			
count	1.304407e+06	1.273286e+06			
mean	2.718484e+02	4.854330e+02			
std	1.175571e+02	1.133489e+02			
min	1.410000e+00	2.160000e+00			
25%	1.828100e+02	4.028600e+02			
50%	2.859400e+02	4.780500e+02			
75%	3.552800e+02	5.597000e+02			
max	7.429600e+02	1.245930e+03			

[8 rows x 22 columns]

```
<ipython-input-6-4ee2834543a3>:4: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
correlation_matrix = df.corr()
```

	Age	BodyweightKg	Squat1Kg	Squat2Kg	Squat3Kg	\
Age	1.000000	0.158509	-0.015515	-0.012730	0.014715	
BodyweightKg	0.158509	1.000000	0.161596	0.126649	0.062187	
Squat1Kg	-0.015515	0.161596	1.000000	0.148796	0.055897	
Squat2Kg	-0.012730	0.126649	0.148796	1.000000	0.128247	
Squat3Kg	0.014715	0.062187	0.055897	0.128247	1.000000	
Squat4Kg	0.054640	0.095291	0.054822	0.096211	0.065949	
Best3SquatKg	0.014667	0.604152	0.332957	0.199968	0.064689	
Bench1Kg	0.025836	0.193994	0.176815	0.127640	0.066851	
Bench2Kg	-0.003216	0.097307	0.110058	0.118657	0.101584	
Bench3Kg	-0.019552	-0.037993	0.044515	0.071562	0.099569	
Bench4Kg	0.069094	0.036419	0.071604	0.045454	0.043731	
Best3BenchKg	0.102148	0.607003	0.277993	0.160325	0.034178	
Deadlift1Kg	0.028104	0.299419	0.190362	0.137626	0.055237	
Deadlift2Kg	0.027182	0.142539	0.087328	0.115457	0.101274	
Deadlift3Kg	0.015448	-0.028371	-0.011092	0.053020	0.115213	
Deadlift4Kg	0.070567	0.018056	-0.036887	0.040110	0.100794	
Best3DeadliftKg	0.027388	0.584668	0.305305	0.190905	0.049729	
TotalKg	-0.136133	0.396248	0.330739	0.201742	0.059239	
Wilks	-0.203642	0.026734	0.269566	0.141056	0.021473	
McCulloch	-0.115789	0.020930	0.246745	0.127062	0.025073	
Glossbrenner	-0.202539	0.045800	0.287243	0.153216	0.024584	
IPFPoints	-0.020378	0.103651	0.252094	0.152462	0.038209	

	Squat4Kg	Best3SquatKg	Bench1Kg	Bench2Kg	Bench3Kg	...	\
Age	0.054640	0.014667	0.025836	-0.003216	-0.019552	...	
BodyweightKg	0.095291	0.604152	0.193994	0.097307	-0.037993	...	
Squat1Kg	0.054822	0.332957	0.176815	0.110058	0.044515	...	
Squat2Kg	0.096211	0.199968	0.127640	0.118657	0.071562	...	
Squat3Kg	0.065949	0.064689	0.066851	0.101584	0.099569	...	
Squat4Kg	1.000000	-0.014048	-0.040992	-0.000900	0.094383	...	
Best3SquatKg	-0.014048	1.000000	0.342576	0.179075	0.015449	...	
Bench1Kg	-0.040992	0.342576	1.000000	0.148200	0.028020	...	
Bench2Kg	-0.000900	0.179075	0.148200	1.000000	0.130057	...	
Bench3Kg	0.094383	0.015449	0.028020	0.130057	1.000000	...	
Bench4Kg	0.297942	0.068681	0.020147	0.072416	0.080381	...	
Best3BenchKg	0.037770	0.884485	0.374812	0.194598	-0.000068	...	
Deadlift1Kg	-0.034457	0.423891	0.230115	0.137560	0.032462	...	
Deadlift2Kg	0.068528	0.126137	0.099606	0.105658	0.073263	...	
Deadlift3Kg	0.087906	-0.151569	-0.025307	0.048406	0.087758	...	
Deadlift4Kg	0.357468	-0.082099	-0.079754	-0.013966	0.077318	...	
Best3DeadliftKg	0.051900	0.888175	0.386443	0.208905	0.021237	...	
TotalKg	0.028978	0.967193	0.194239	0.131773	0.058817	...	
Wilks	-0.028537	0.774673	0.081439	0.076524	0.071695	...	
McCulloch	-0.011682	0.740720	0.070058	0.071619	0.074931	...	
Glossbrenner	-0.026084	0.817954	0.097963	0.084838	0.070113	...	
IPFPoints	-0.053050	0.632758	0.251453	0.148762	0.022306	...	

	Deadlift1Kg	Deadlift2Kg	Deadlift3Kg	Deadlift4Kg	\
Age	0.028104	0.027182	0.015448	0.070567	
BodyweightKg	0.299419	0.142539	-0.028371	0.018056	
Squat1Kg	0.190362	0.087328	-0.011092	-0.036887	
Squat2Kg	0.137626	0.115457	0.053020	0.040110	
Squat3Kg	0.055237	0.101274	0.115213	0.100794	
Squat4Kg	-0.034457	0.068528	0.087906	0.357468	
Best3SquatKg	0.423891	0.126137	-0.151569	-0.082099	
Bench1Kg	0.230115	0.099606	-0.025307	-0.079754	
Bench2Kg	0.137560	0.105658	0.048406	-0.013966	
Bench3Kg	0.032462	0.073263	0.087758	0.077318	



Bench4Kg	-0.011695	0.015117	0.057702	0.347674
Best3BenchKg	0.413953	0.137041	-0.120944	-0.026343
Deadlift1Kg	1.000000	0.115412	-0.070460	-0.083301
Deadlift2Kg	0.115412	1.000000	0.141795	-0.037798
Deadlift3Kg	-0.070460	0.141795	1.000000	0.106620
Deadlift4Kg	-0.083301	-0.037798	0.106620	1.000000
Best3DeadliftKg	0.531479	0.221535	-0.076600	-0.090372
TotalKg	0.453955	0.162569	-0.100441	-0.064316
Wilks	0.307736	0.091039	-0.105416	-0.076961
McCulloch	0.283836	0.085846	-0.096568	-0.057145
Glossbrenner	0.338074	0.103341	-0.109219	-0.076339
IPFPoints	0.337982	0.122884	-0.090485	-0.156644

	Best3DeadliftKg	TotalKg	Wilks	McCulloch	Glossbrenner	\
Age	0.027388	-0.136133	-0.203642	-0.115789	-0.202539	
BodyweightKg	0.584668	0.396248	0.026734	0.020930	0.045800	
Squat1Kg	0.305305	0.330739	0.269566	0.246745	0.287243	
Squat2Kg	0.190905	0.201742	0.141056	0.127062	0.153216	
Squat3Kg	0.049729	0.059239	0.021473	0.025073	0.024584	
Squat4Kg	0.051900	0.028978	-0.028537	-0.011682	-0.026084	
Best3SquatKg	0.888175	0.967193	0.774673	0.740720	0.817954	
Bench1Kg	0.386443	0.194239	0.081439	0.070058	0.097963	
Bench2Kg	0.208905	0.131773	0.076524	0.071619	0.084838	
Bench3Kg	0.021237	0.058817	0.071695	0.074931	0.070113	
Bench4Kg	0.092414	0.025473	0.005965	0.011905	0.008602	
Best3BenchKg	0.866538	0.483902	0.184129	0.173727	0.227361	
Deadlift1Kg	0.531479	0.453955	0.307736	0.283836	0.338074	
Deadlift2Kg	0.221535	0.162569	0.091039	0.085846	0.103341	
Deadlift3Kg	-0.076600	-0.100441	-0.105416	-0.096568	-0.109219	
Deadlift4Kg	-0.090372	-0.064316	-0.076961	-0.057145	-0.076339	
Best3DeadliftKg	1.000000	0.864799	0.607853	0.588145	0.661566	
TotalKg	0.864799	1.000000	0.881385	0.867507	0.906568	
Wilks	0.607853	0.881385	1.000000	0.985428	0.995393	
McCulloch	0.588145	0.867507	0.985428	1.000000	0.981328	
Glossbrenner	0.661566	0.906568	0.995393	0.981328	1.000000	
IPFPoints	0.673613	0.267329	0.263062	0.245303	0.262000	

	IPFPoints
Age	-0.020378
BodyweightKg	0.103651
Squat1Kg	0.252094
Squat2Kg	0.152462
Squat3Kg	0.038209
Squat4Kg	-0.053050
Best3SquatKg	0.632758
Bench1Kg	0.251453
Bench2Kg	0.148762
Bench3Kg	0.022306
Bench4Kg	-0.112325
Best3BenchKg	0.621578
Deadlift1Kg	0.337982
Deadlift2Kg	0.122884
Deadlift3Kg	-0.090485
Deadlift4Kg	-0.156644
Best3DeadliftKg	0.673613
TotalKg	0.267329
Wilks	0.263062
McCulloch	0.245303
Glossbrenner	0.262000

IPFPoints 1.000000

[22 rows x 22 columns]

Missing values count: Name 0

Sex 0  
Event 0  
Equipment 0  
Age 665827  
AgeClass 636554  
Division 8178  
BodyweightKg 16732  
WeightClassKg 13312  
Squat1Kg 1085774  
Squat2Kg 1090005  
Squat3Kg 1099512  
Squat4Kg 1419658  
Best3SquatKg 391904  
Bench1Kg 923575  
Bench2Kg 929868  
Bench3Kg 944869  
Bench4Kg 1413849  
Best3BenchKg 147173  
Deadlift1Kg 1059810  
Deadlift2Kg 1067331  
Deadlift3Kg 1083407  
Deadlift4Kg 1414108  
Best3DeadliftKg 341546  
TotalKg 110170  
Place 0  
Wilks 118947  
McCulloch 119100  
Glossbrenner 118947  
IPFPoints 150068  
Tested 329462  
Country 1034470  
Federation 0  
Date 0  
MeetCountry 0  
MeetState 481809  
MeetName 0

dtype: int64

Total NaN values: 18215965

Total duplicated rows count: 3084

	Name	Sex	Event	Equipment	Age	AgeClass	Division	BodyweightKg	\
0	Abbie Murphy	F	SBD	Wraps	29.0	24-34	F-OR	59.8	
1	Abbie Tuong	F	SBD	Wraps	29.0	24-34	F-OR	58.5	
2	Ainslee Hooper	F	B	Raw	40.0	40-44	F-OR	55.4	
3	Amy Moldenhauer	F	SBD	Wraps	23.0	20-23	F-OR	60.0	
4	Andrea Rowan	F	SBD	Wraps	45.0	45-49	F-OR	104.0	
	WeightClassKg	Squat1Kg	...	McCulloch	Glossbrenner	IPFPoints	Tested	\	
0	60	80.0	...	324.16	286.42	511.15	NaN		
1	60	100.0	...	378.07	334.16	595.65	NaN		
2	56	NaN	...	38.56	34.12	313.97	NaN		
3	60	-105.0	...	345.61	305.37	547.04	NaN		
4	110	120.0	...	338.91	274.56	550.08	NaN		
	Country	Federation	Date	MeetCountry	MeetState	MeetName			
0	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup			

1	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
2	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
3	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup
4	NaN	GPC-AUS	2018-10-27	Australia	VIC	Melbourne Cup

[5 rows x 37 columns]

As we can see from the graphs a large majority of lifters are male, typically powerlifting is a male dominated sport, but it has had a rise in popularity and more women are competing. Most of the male and female lifters are also between the ages of 15 and 45, with the highest being between ages of 20 and 25, so this is more likely where we will see the strongest athletes. The majority of male lifters are between 75 and 90 kilograms. The majority of female lifters are between 50 and 60 kilograms.

## Section 2a.

In this section we use our primary data cleaning technique

```
In [ ]: from sklearn.preprocessing import LabelEncoder
# Some columns aren't really necessary so we drop the irrelevant columns
df = df.drop(['Squat1Kg', 'Squat2Kg', 'Squat3Kg', 'Squat4Kg',
              'Best3SquatKg', 'Bench1Kg', 'Bench2Kg', 'Bench3Kg', 'Bench4Kg',
              'Best3BenchKg', 'Deadlift1Kg', 'Deadlift2Kg', 'Deadlift3Kg', 'Deadlift4Kg',
              'Best3DeadliftKg', 'AgeClass', 'Federation', 'Tested', 'Date', 'MeetCountry',
              'MeetState', 'MeetName', 'Place', 'Country', 'Wilks', 'McCulloch',
              'Glossbrenner', 'IPFPoints', 'Event'], axis=1)
# We want to filter the dataset to include only raw and wraps equipment lifters
df_raw_wraps = copy.deepcopy(df)
df_raw_wraps = df_raw_wraps.loc[(df_raw_wraps['Equipment'] == 'Wraps') | (df_raw_wraps['Equipment'] == 'Raw')]
# There are various missing values as we can see so we drop them also we drop any values that are not numeric and we drop duplicates
df_raw_wraps.loc[:, 'Age'] = pd.to_numeric(df_raw_wraps['Age'], errors='coerce')
df_raw_wraps.loc[:, 'BodyweightKg'] = pd.to_numeric(df_raw_wraps['BodyweightKg'], errors='coerce')
df_raw_wraps.loc[:, 'WeightClassKg'] = pd.to_numeric(df_raw_wraps['WeightClassKg'], errors='coerce')
df_raw_wraps.loc[:, 'TotalKg'] = pd.to_numeric(df_raw_wraps['TotalKg'], errors='coerce')
df_raw_wraps = df_raw_wraps.drop_duplicates()
features = ['Sex', 'Age', 'WeightClassKg', 'BodyweightKg', 'Division']
df_raw_wraps = df_raw_wraps.dropna(subset=features)
df_raw_wraps.dropna(subset=['TotalKg'], inplace=True)
```

```
<ipython-input-7-ee73549aeeb8>:13: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_raw_wraps.loc[:, 'Age'] = pd.to_numeric(df_raw_wraps['Age'], errors='coerce')
<ipython-input-7-ee73549aeeb8>:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_raw_wraps.loc[:, 'BodyweightKg'] = pd.to_numeric(df_raw_wraps['BodyweightKg'], errors='coerce')
<ipython-input-7-ee73549aeeb8>:15: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values inplace instead of always setting a new array.
To retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns are non-unique, `df.isetitem(i, newvals)`
df_raw_wraps.loc[:, 'WeightClassKg'] = pd.to_numeric(df_raw_wraps['WeightClassKg'], errors='coerce')
<ipython-input-7-ee73549aeeb8>:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_raw_wraps.dropna(subset=['TotalKg'], inplace=True)
```

## Section 2b.

In this section we use our primary data processing technique

```
In [ ]: # Here we encode categorical columns to numerical values, we used this during lecture so I will use for this data cleaning
le = LabelEncoder()
df_raw_wraps['Sex'] = le.fit_transform(df_raw_wraps['Sex'])
df_raw_wraps['Equipment'] = le.fit_transform(df_raw_wraps['Equipment'])
df_raw_wraps['Division'] = le.fit_transform(df_raw_wraps['Division'])
# In this section we are normalizing the features using sklearn, and defining X and y, X refers to our feature data that we will use to predict our target variable y
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaleable_features = ['Age', 'WeightClassKg', 'BodyweightKg']
df_raw_wraps[scaleable_features] = scaler.fit_transform(df_raw_wraps[scaleable_features])
X = df_raw_wraps[scaleable_features]
y = df_raw_wraps['TotalKg']
# We make a copy of the dataframe since we want to unscale the values as we will need to use this for data analysis
df_unscaled = df_raw_wraps.copy()
df_unscaled[scaleable_features] = scaler.inverse_transform(df_raw_wraps[scaleable_features])
```

```
<ipython-input-8-808cf907411f>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_raw_wraps['Sex'] = le.fit_transform(df_raw_wraps['Sex'])
```

## Section 3a.

In this section we utilize a different data cleaning technique

```
In [ ]: # Instead of dropping columns we just keep the columns of our features and totalkg which is what we are trying to predict
df_copy = df_copy[['Sex', 'Equipment', 'Age', 'WeightClassKg', 'BodyweightKg', 'TotalKg']]
df_copy = df_copy.dropna(subset=['TotalKg'])
df_copy = df_copy.drop_duplicates()
df_copy = df_copy[df_copy['Equipment'].isin(['Raw', 'Wraps'])]
df_copy.reset_index(drop=True, inplace=True)
temp = ['Age', 'BodyweightKg', 'TotalKg']
for i in range(len(temp)):
    df_copy[temp[i]] = pd.to_numeric(df_copy[temp[i]], errors='coerce')
df_copy['WeightClassKg'] = df_copy['WeightClassKg'].str.replace('+', '')
df_copy['WeightClassKg'] = pd.to_numeric(df_copy['WeightClassKg'], errors='coerce')
```

```
<ipython-input-9-882e637e4d89>:10: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.
df_copy['WeightClassKg'] = df_copy['WeightClassKg'].str.replace('+', '')
```

## Section 3b.

In this section we use our secondary data processing technique

```
In [ ]: from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler
# We will utilize one-hot encoding for categorical features
OHE = OneHotEncoder(sparse=False)
Encoded_Sex = OHE.fit_transform(df_copy[['Sex']])
sex_columns = OHE.categories_[0].tolist()
df_raw_wraps_encoded = pd.DataFrame(Encoded_Sex, columns=sex_columns, index=df_copy.index)
Encoded_Equipment = OHE.fit_transform(df_copy[['Equipment']])
equipment_columns = OHE.categories_[0].tolist()
df_raw_wraps_encoded = pd.concat([df_raw_wraps_encoded, pd.DataFrame(Encoded_Equipment, columns=equipment_columns, index=df_copy.index)], axis=1)
# We have many missing values in the dataset on way to handle the missing values is to impute the missing values using the mean of each column
im = SimpleImputer(strategy='mean')
temp = df_copy[['Age', 'WeightClassKg', 'BodyweightKg']]
imputed_df = pd.DataFrame(im.fit_transform(temp), columns=temp.columns, index=temp.index)
df_raw_wraps_encoded = pd.concat([df_raw_wraps_encoded, imputed_df], axis=1)

# Here we normalize age, weight class, and body weight
MMS = MinMaxScaler()
scaled_data = MMS.fit_transform(df_raw_wraps_encoded[['Age', 'WeightClassKg', 'BodyweightKg']])
df_raw_wraps_encoded[['Age', 'WeightClassKg', 'BodyweightKg']] = scaled_data

X_prime = df_raw_wraps_encoded
y_prime = df_copy['TotalKg']
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(
```

## Section 4.

In this section we split the dataset into training and testing

```
In [ ]: # In this section we are splitting the data into training and testing for both techniques used
# a common test split we discussed is 10% testing with 10% validation and 80% training the random state is set to 7
from sklearn.model_selection import train_test_split
#Here we split the data of our initial data processing and cleaning into training, testing and validation.
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.2, random_state=7)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=.5, random_state=7)
# Here we split the data of our alternative data processing and cleaning into training, testing and validation.
X_prime_train, X_prime_temp, y_prime_train, y_prime_temp = train_test_split(X_prime, y_prime, test_size=0.3, random_state=7)
X_prime_val, X_prime_test, y_prime_val, y_prime_test = train_test_split(X_prime_temp, y_prime_temp, test_size=0.33, random_state=7)

X_unscaled = df_unscaled[features]
y_unscaled = df_unscaled['TotalKg']
X_train_unscaled, X_test_unscaled, y_train_unscaled, y_test_unscaled = train_test_split(X_unscaled, y_unscaled, test_size=0.2, random_state=7)
```

## Section 5a.

In this section we perform data analysis on the male training subset

```
In [ ]: # In this section we look at the descriptive statistics of the male competitors
male_df = X_train_unscaled[X_train_unscaled['Sex'] == 1]
print('-----')
print("Descriptive statistics of male competitors:")
print(male_df[['Age', 'BodyweightKg']].describe())
print('-----')
plt.figure(figsize=(20, 10))
sns.histplot(data=male_df, x='Age', kde=True, label='Male')
plt.legend()
plt.title('Distribution of Age for Male Competitors')
plt.show()

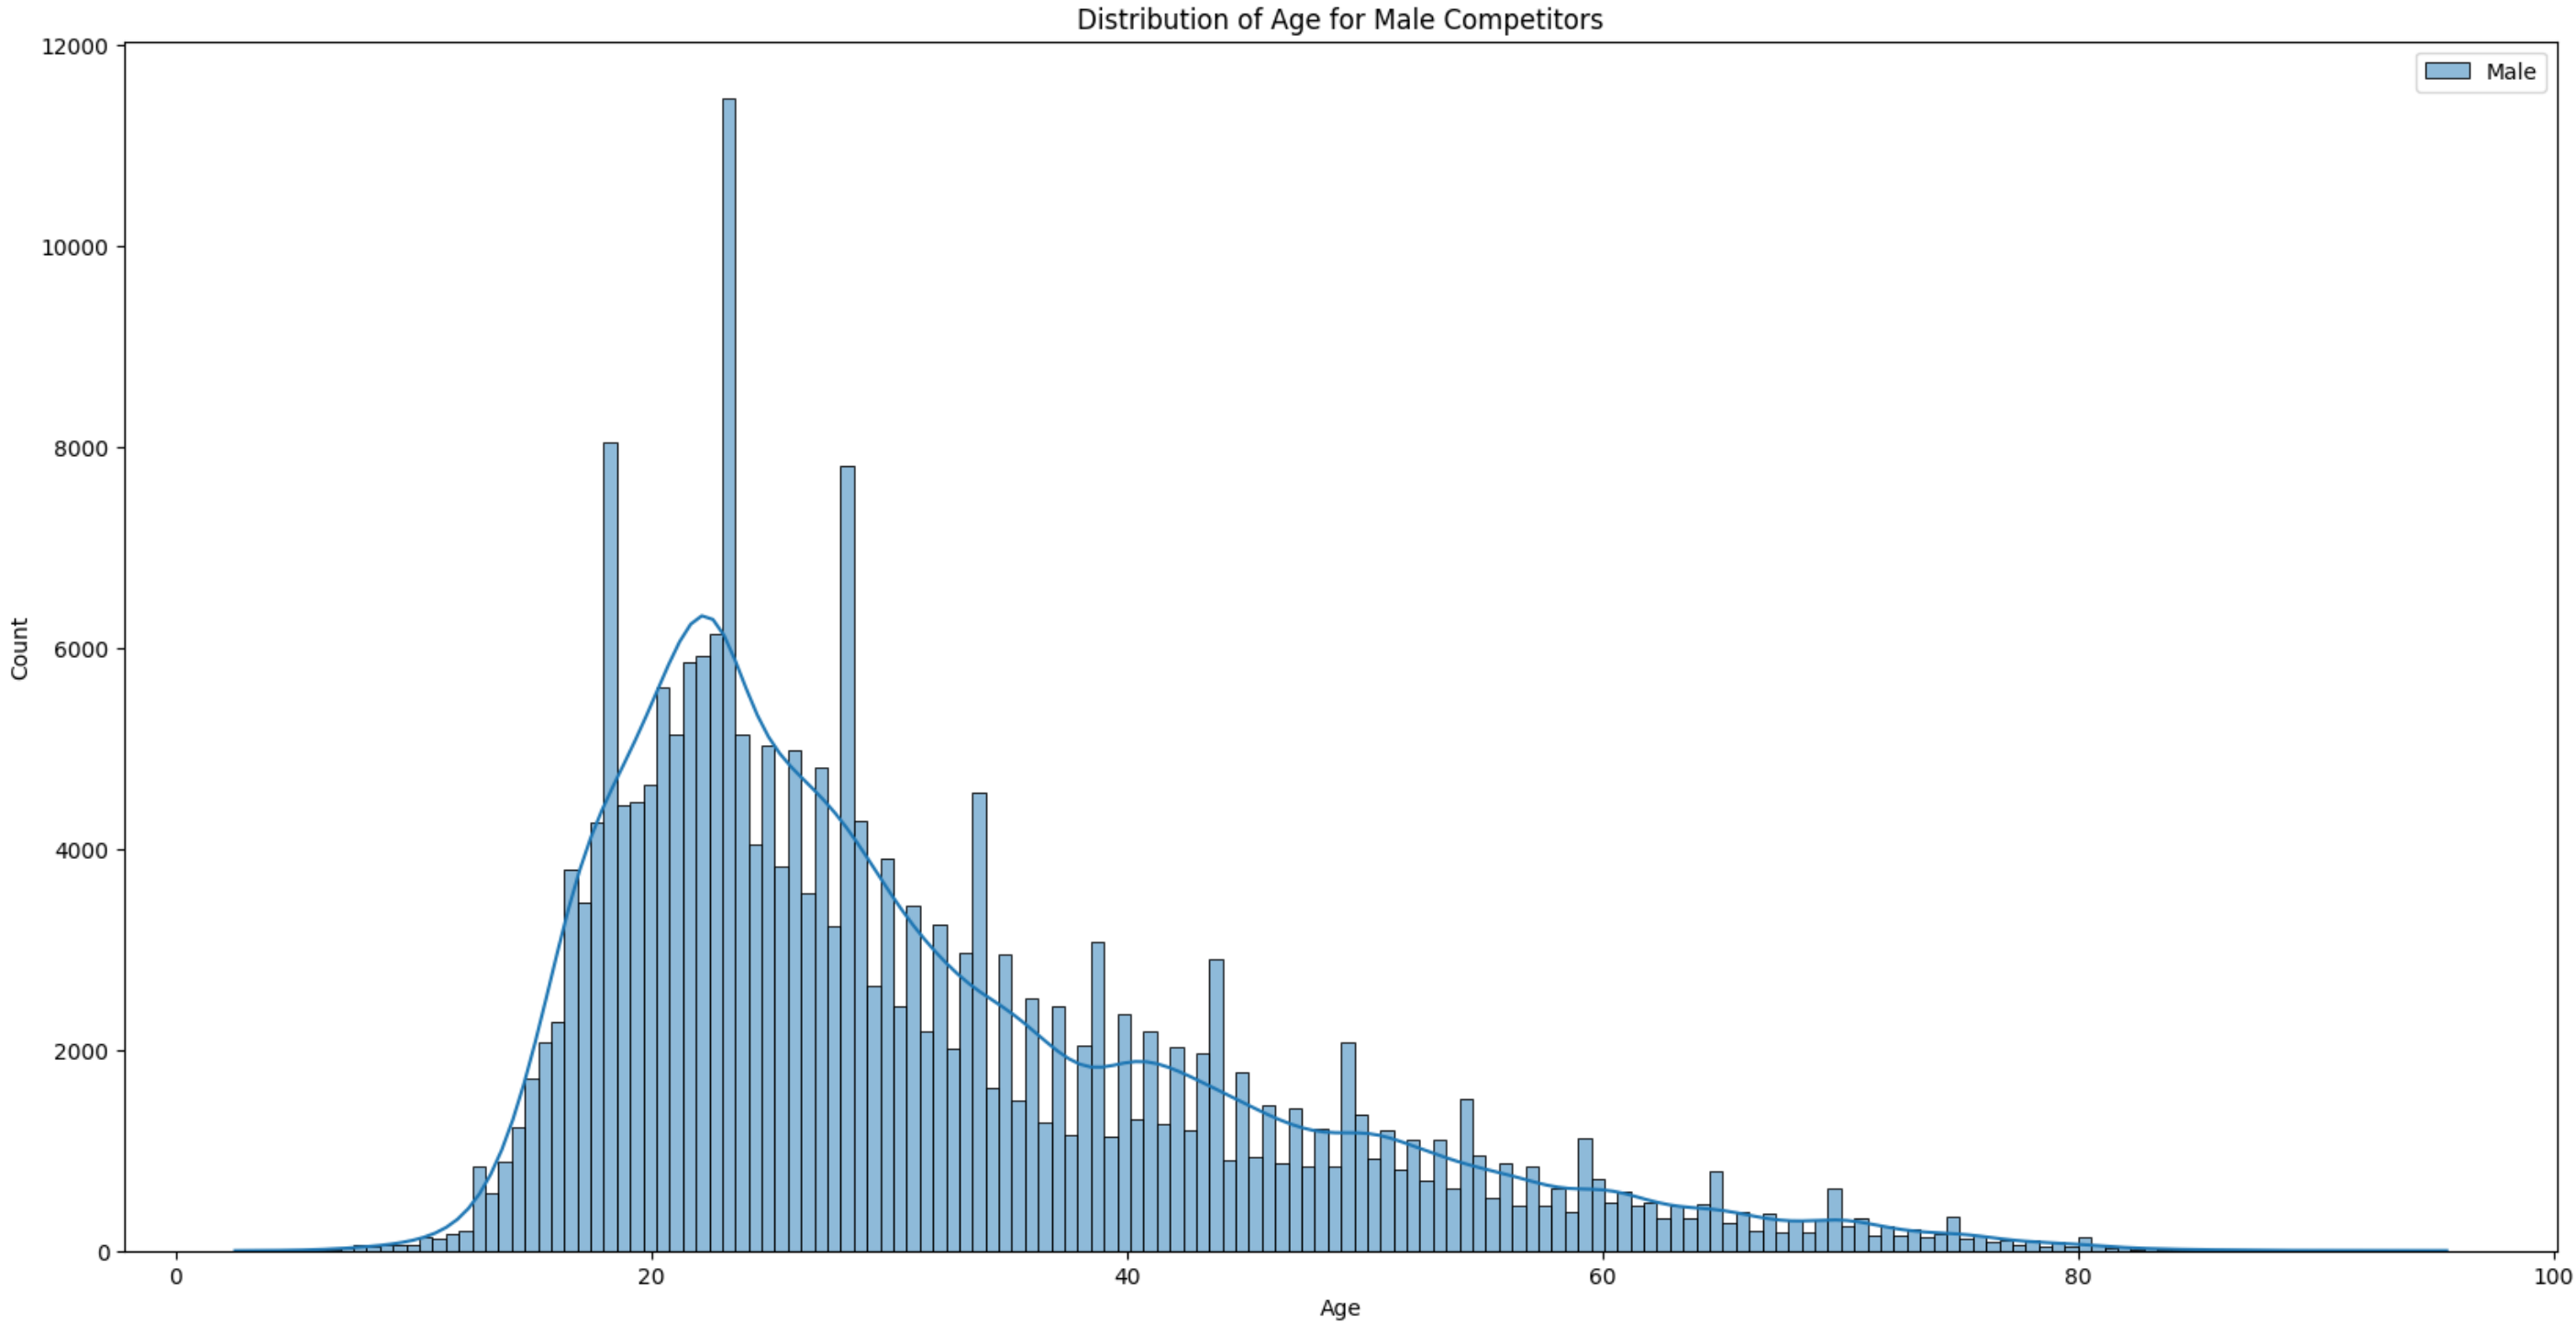
plt.figure(figsize=(20, 10))
sns.histplot(data=male_df, x='BodyweightKg', kde=True, label='Male')
plt.legend()
plt.title('Distribution of Bodyweight for Male Competitors')
plt.show()
```


-----

Descriptive statistics of male competitors:

	Age	BodyweightKg
count	231672.000000	231672.000000
mean	31.354441	89.737960
std	13.405141	18.080501
min	2.500000	17.690000
25%	21.500000	76.500000
50%	27.500000	88.900000
75%	38.500000	101.500000
max	95.500000	155.000000

-----





## Section 5b.

In this section we perform data analysis on the female training subset



```
In [ ]: # In this section we look at the despritive statistics of the female competitors
female_df = X_train_unscaled[X_train_unscaled['Sex'] == 0]
print('-----')
print("Descriptive statistics of female competitors:")
print(female_df[['Age', 'BodyweightKg']].describe())
print('-----')
plt.figure(figsize=(20, 10))
sns.histplot(data=female_df, x='Age', kde=True, label='Female')
plt.legend()
plt.title('Distribution of Age for Female Competitors')
plt.show()

plt.figure(figsize=(20, 10))
sns.histplot(data=female_df, x='BodyweightKg', kde=True, label='Female')
plt.legend()
plt.title('Distribution of Bodyweight for Female Competitors')
plt.show()
```


-----


Descriptive statistics of female competitors:

	Age	BodyweightKg
count	84788.000000	84788.000000
mean	31.289215	64.816118
std	12.083642	11.200386
min	0.500000	15.880000
25%	22.500000	56.200000
50%	28.500000	63.800000
75%	38.500000	72.050000
max	97.000000	125.000000

-----

Distribution of Age for Female Competitors





Sections 7a-7b is used for fine tuning the models, the performance of the model is evaluated on the validation set. We define a function that allows us to evaluate model performance.

```
In [ ]: def evaluate_model(model, X_train, y_train, X_val, y_val, title):
        model.fit(X_train, y_train)
        y_pred_val = model.predict(X_val)


        plt.figure(figsize=(10, 5))
        plt.scatter(y_val, y_pred_val, alpha=0.5)
        plt.plot([min(y_val), max(y_val)], [min(y_val), max(y_val)], color='red')
        plt.xlabel('Actual values')
        plt.ylabel('Predicted values')
        plt.title(title)
        plt.show()

        mae = mean_absolute_error(y_val, y_pred_val)
        mse = mean_squared_error(y_val, y_pred_val)
        r2 = r2_score(y_val, y_pred_val)
        print("Mean Absolute Error: " + str(mae))
        print("Mean Squared Error: " + str(mse))
        print("R-squared: " + str(r2))
```


## Section 6.

In this section we implement a base model

```
In [ ]: from sklearn.linear_model import LinearRegression
        from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
        # We implement linear regression and print the models performance
        lr = LinearRegression()
        evaluate_model(lr, X_train, y_train, X_val, y_val, 'Linear regression model performance')
        evaluate_model(lr, X_prime_train, y_prime_train, X_prime_val, y_prime_val, 'Linear regression model performance')
```



Mean Absolute Error: 149.03692628018624  
Mean Squared Error: 31114.539205523048  
R-squared: 0.24317267369691398




Mean Absolute Error: 140.8475128547402  
Mean Squared Error: 27620.004192423316  
R-squared: 0.33162761874234414


## Section 7a.

In this section we implement KNN

```
In [ ]: from sklearn.neighbors import KNeighborsRegressor
# We implement KKN alorithm we start with k = 5
knn = KNeighborsRegressor(n_neighbors=5)
evaluate_model(knn, X_train, y_train, X_val, y_val, 'KNN model performance')
evaluate_model(knn, X_prime_train, y_prime_train, X_prime_val, y_prime_val, 'KNN model performance')
```



Mean Absolute Error: 129.9284232373537  
Mean Squared Error: 28735.49088119251  
R-squared: 0.3010404367563557




Mean Absolute Error: 122.56268874051409  
Mean Squared Error: 25427.84724494765  
R-squared: 0.3846752992881982

## Section 7b.

In this section we implement MLP


```
In [ ]: from sklearn.neural_network import MLPRegressor
# We implement MLP the activation is relu, with 100 hidden layers, random state is as usual 7.
mlp = MLPRegressor(
    activation='relu',
    hidden_layer_sizes=(10, 100),
    alpha=0.001,
    random_state=7,
    early_stopping=False
)
evaluate_model(mlp, X_train, y_train, X_val, y_val, 'MLP model performance')
evaluate_model(mlp, X_prime_train, y_prime_train, X_prime_val, y_prime_val, 'MLP model performance')
```





Mean Absolute Error: 135.41043254080915  
Mean Squared Error: 28017.956156034612  
R-squared: 0.3184936885620092

```
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.  
  warnings.warn(
```



Mean Absolute Error: 129.89485201812323  
Mean Squared Error: 24898.968521083167  
R-squared: 0.3974735570148493

## Section 8.

In this section we look at training performance and validation performance. We define a function `model_performance`, this allows us to see if the model overfit or underfit.

```
In [ ]: def model_performance(model, X_train, y_train, X_val, y_val, model_name):
    model.fit(X_train, y_train)

    y_pred_train = model.predict(X_train)
    y_pred_val = model.predict(X_val)

    mae_train = mean_absolute_error(y_train, y_pred_train)
    mse_train = mean_squared_error(y_train, y_pred_train)
    r2_train = r2_score(y_train, y_pred_train)

    mae_val = mean_absolute_error(y_val, y_pred_val)
    mse_val = mean_squared_error(y_val, y_pred_val)
    r2_val = r2_score(y_val, y_pred_val)
    print("Model: " + model_name)
    print("Training performance:")
    print("Mean Absolute Error: " + str(mae_train))
    print("Mean Squared Error: " + str(mse_train))
    print("R-squared: " + str(r2_train))
    print("Validation performance:")
    print("Mean Absolute Error: " + str(mae_val))
    print("Mean Squared Error: " + str(mse_val))
    print("R-squared: " + str(r2_val))
```

```
In [ ]: lr = LinearRegression()
knn = KNeighborsRegressor(n_neighbors=5)
mlp = MLPRegressor(
    activation='relu',
    hidden_layer_sizes=(10, 100),
    alpha=0.001,
    random_state=7,
    early_stopping=False
)
model_performance(lr, X_train, y_train, X_val, y_val, "Linear regression original data cleaning and processing")
model_performance(lr, X_prime_train, y_prime_train, X_prime_val, y_prime_val, "Linear regression alternative data cleaning and processing")

model_performance(knn, X_train, y_train, X_val, y_val, "K Nearest Neighbor original data cleaning and processing")
model_performance(knn, X_prime_train, y_prime_train, X_prime_val, y_prime_val, "K Nearest Neighbor alternative data cleaning and processing")

model_performance(mlp, X_train, y_train, X_val, y_val, "Mutilayer perceptron original data cleaning and processing")
model_performance(mlp, X_prime_train, y_prime_train, X_prime_val, y_prime_val, "Mutilayer perceptron alternative data cleaning and processing")
```

```
Model: Linear regression original data cleaning and proccesing
Training performance:
Mean Absolute Error: 149.27090152337968
Mean Squared Error: 31157.59819147712
R-squared: 0.2475899161253794
Validation performance:
Mean Absolute Error: 149.03692628018624
Mean Squared Error: 31114.539205523048
R-squared: 0.2475899161253794
Model: Linear regression alternative data cleaning and proccesing
Training performance:
Mean Absolute Error: 141.2584858062405
Mean Squared Error: 27785.409727887098
R-squared: 0.32870425310137497
Validation performance:
Mean Absolute Error: 140.8475128547402
Mean Squared Error: 27620.004192423316
R-squared: 0.32870425310137497
Model: K Nearest Neighbor original data cleaning and proccesing
Training performance:
Mean Absolute Error: 106.12909671364473
Mean Squared Error: 19399.78974995841
R-squared: 0.531523664204997
Validation performance:
Mean Absolute Error: 129.9284232373537
Mean Squared Error: 28735.49088119251
R-squared: 0.531523664204997
Model: K Nearest Neighbor alternative data cleaning and proccesing
Training performance:
Mean Absolute Error: 103.56414726889227
Mean Squared Error: 18490.29518895408
R-squared: 0.5532743032834568
Validation performance:
Mean Absolute Error: 122.56268874051409
Mean Squared Error: 25427.84724494765
R-squared: 0.5532743032834568
Model: Mutilayer perceptron original data cleaning and proccesing
Training performance:
Mean Absolute Error: 135.3840051940262
Mean Squared Error: 28020.113647364185
R-squared: 0.32335554460816407
Validation performance:
Mean Absolute Error: 135.41043254080915
Mean Squared Error: 28017.956156034612
R-squared: 0.32335554460816407
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the op
timization hasn't converged yet.
```

```
warnings.warn(
```

```
Model: Mutilayer perceptron alternative data cleaning and proccesing
Training performance:
Mean Absolute Error: 130.0991947118325
Mean Squared Error: 25014.240899441553
R-squared: 0.39565571671814004
Validation performance:
Mean Absolute Error: 129.89485201812323
Mean Squared Error: 24898.968521083167
R-squared: 0.39565571671814004
```

## Section 9.

In this section we train and evaluate the performance of the model on the test set. We also time the each model and obtain the memory size of the model. We define a function `train_and_evaluate`, which returns `total_training_time`, `memory_size`, `mae_test`, `mse_test`, `r2_test`.

```
In [ ]: import time
import sys

def train_and_evaluate(model, X_train, y_train, X_test, y_test,model_name):
    start = time.time()
    model.fit(X_train, y_train)
    end = time.time()
    total_training_time = end - start
    memory_size = sys.getsizeof(model)

    y_pred_test = model.predict(X_test)
    mae_test = mean_absolute_error(y_test, y_pred_test)
    mse_test = mean_squared_error(y_test, y_pred_test)
    r2_test = r2_score(y_test, y_pred_test)

    return [model_name,total_training_time, memory_size, mae_test, mse_test, r2_test]

def display_results(Original, Alternative):
    print("Model:" +str(Original[0]))
    columns = ['model_name', 'Training Time', 'Memory Size', 'MAE', 'MSE', 'R^2']
    df = pd.DataFrame([Original, Alternative], columns=columns)
    for i in columns[1:]:
        print("Attribute : " +str(i))
        ax = df[i].plot(kind='bar', figsize=(6, 4), color=['blue', 'orange'])
        ax.set_ylabel('Value')
        ax.set_title(f'Comparison of {i} between original and alternative data cleaning and processing')
        ax.legend(['Original', 'Alternative'])
        plt.tight_layout()
        plt.show()

def print_results(results):
    print('-----')
    print("Model name: " + str(results[0]))
    print("Total training time : " + str(results[1]) + " seconds")
    print("Memory size: " + str(results[2]) + " bytes")
    print("Mean Absolute Error : " + str(results[3]))
    print("Mean Squared Error : " + str(results[4]))
    print("R-squared: " + str(results[5]))
    print('-----')
```

```
In [ ]: lr = LinearRegression()
Original_lr = train_and_evaluate(lr, X_train, y_train, X_test, y_test,"linear regression")
Alternative_lr = train_and_evaluate(lr, X_prime_train, y_prime_train, X_prime_test, y_prime_test,"linear regression")
```

```
In [ ]: knn = KNeighborsRegressor(n_neighbors=7)
Original_knn= train_and_evaluate(knn, X_train, y_train, X_test, y_test, " k-nearest-neighbor")
Alternative_knn = train_and_evaluate(knn, X_prime_train, y_prime_train, X_prime_test, y_prime_test, " k-nearest-neighbor")
```

```
In [ ]: mlp = MLPRegressor(
    activation='relu',
    hidden_layer_sizes=(10, 100),
    alpha=0.001,
    random_state=7,
    early_stopping=False
)
Original_mlp = train_and_evaluate(mlp, X_train, y_train, X_test, y_test, "multilayer perceptron")
Alternative_mlp = train_and_evaluate(mlp, X_prime_train, y_prime_train, X_prime_test, y_prime_test, "multilayer perceptron")
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
  warnings.warn(
```

## Section 10.

In this section we print and display results, print\_results shows us numerically the results, display\_results shows us visually

```
In [ ]: print_results(Original_lr)
print_results(Alternative_lr)
print_results(Original_knn)
print_results(Alternative_knn)
print_results(Original_mlp)
print_results(Alternative_mlp)
```

```

-----
Model name: linear regression
Total training time : 0.07825016975402832 seconds
Memory size: 48 bytes
Mean Absolute Error : 148.56010580663022
Mean Squared Error : 30980.337804732517
R-squared: 0.2505673863799529
-----
Model name: linear regression
Total training time : 0.10439372062683105 seconds
Memory size: 48 bytes
Mean Absolute Error : 140.77681677226255
Mean Squared Error : 27644.238058774758
R-squared: 0.33121514168764
-----
Model name: k-nearest-neighbor
Total training time : 0.46447086334228516 seconds
Memory size: 48 bytes
Mean Absolute Error : 127.8504918275516
Mean Squared Error : 27358.322241304737
R-squared: 0.3381860756073347
-----
Model name: k-nearest-neighbor
Total training time : 0.5056524276733398 seconds
Memory size: 48 bytes
Mean Absolute Error : 119.7086877995598
Mean Squared Error : 24122.29349400316
R-squared: 0.41641999311913425
-----
Model name: multilayer perceptron
Total training time : 306.41275668144226 seconds
Memory size: 48 bytes
Mean Absolute Error : 135.04628374208875
Mean Squared Error : 27910.389030681305
R-squared: 0.32483125489934017
-----
Model name: multilayer perceptron
Total training time : 1124.6327872276306 seconds
Memory size: 48 bytes
Mean Absolute Error : 129.39169318843508
Mean Squared Error : 24773.67200870933
R-squared: 0.4006614800163575
-----

```

```

In [ ]: display_results(Original_lr, Alternative_lr)
display_results(Original_knn, Alternative_knn)
display_results(Original_mlp, Alternative_mlp)


```

```

Model:linear regression
Attribute : Training Time


```

Comparison of Training Time between original and alternative data cleaning and processing



Attribute : Memory Size


Comparison of Memory Size between original and alternative data cleaning and processing



Attribute : MAE




Comparison of MAE between original and alternative data cleaning and processing




Attribute : MSE

Comparison of MSE between original and alternative data cleaning and processing




Attribute : R^2

Comparison of R^2 between original and alternative data cleaning and processing




Model: k-nearest-neighbor  
Attribute : Training Time

Comparison of Training Time between original and alternative data cleaning and processing




Attribute : Memory Size

Comparison of Memory Size between original and alternative data cleaning and processing




Attribute : MAE

Comparison of MAE between original and alternative data cleaning and processing




Attribute : MSE

Comparison of MSE between original and alternative data cleaning and processing




Attribute : R^2

Comparison of R^2 between original and alternative data cleaning and processing




Model:multilayer perceptron  
Attribute : Training Time

Comparison of Training Time between original and alternative data cleaning and processing




Attribute : Memory Size

Comparison of Memory Size between original and alternative data cleaning and processing




Attribute : MAE

Comparison of MAE between original and alternative data cleaning and processing



Attribute : MSE

Comparison of MSE between original and alternative data cleaning and processing



Attribute : R^2

Comparison of R^2 between original and alternative data cleaning and processing

